

# SeaHorn

A fully automated analysis framework  
for LLVM-based languages

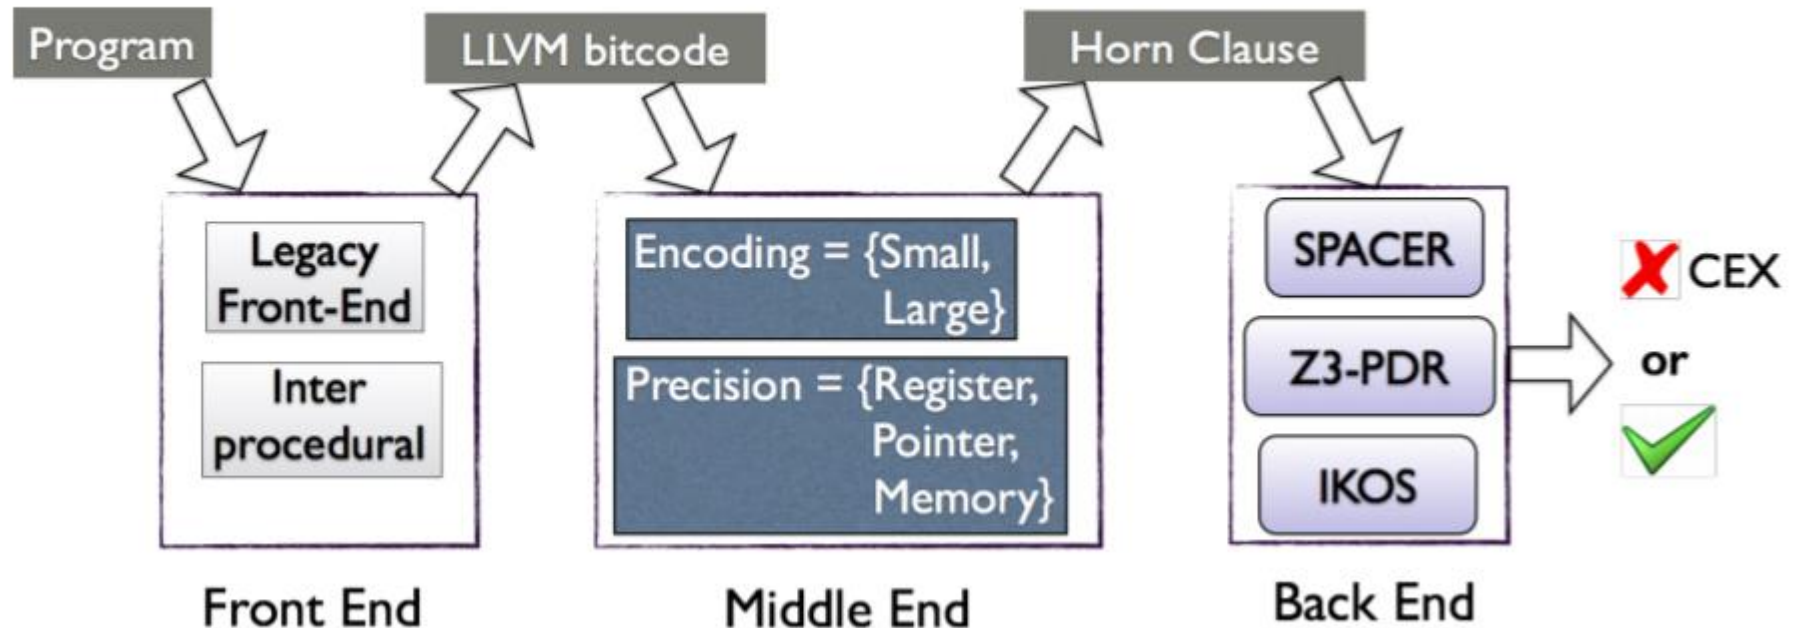
Weizhi Feng

January 12, 2021

# Overview

- SeaHorn is an automated analysis framework for LLVM-based languages.

## SeaHorn Verification Framework

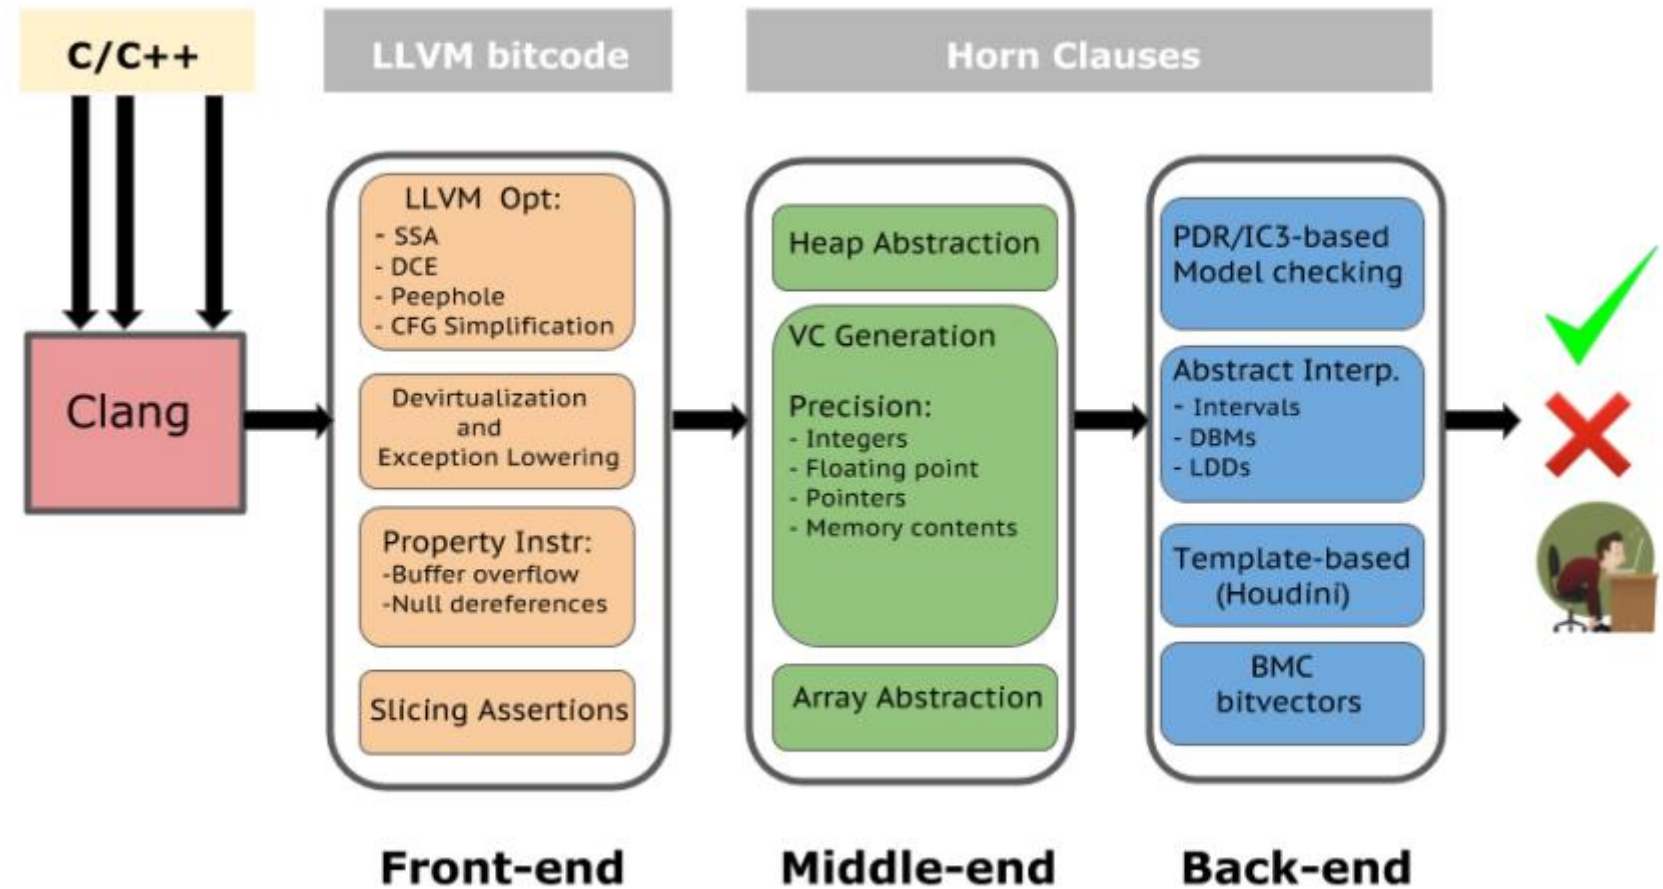


# Overview

- ▶ SeaHorn is an automated analysis framework for LLVM-based languages.
- ▶ Distinguishing Features
  - LLVM front-end
  - Constrained horn clauses to represent verification conditions
  - Comparable to state-of-the-art tools at SV-COMP 15
- ▶ Usage
  - > sea pf FILE.c
    - ✓ Outputs sat for unsafe (has counterexample); unsat for safe.
  - Additional options
    - ✓ --cex=trace.xml outputs a counter-example in SV-COMP 15 format.
    - ✓ --track={reg,ptr,mem} track registers, points, memory content.
    - ✓ --step={large, small} verification condition step-semantics,
    - ✓ -small == basic block, large == loop-free control flow block.
    - ✓ --inline inline all functions in the front-end passes.

# Workflow

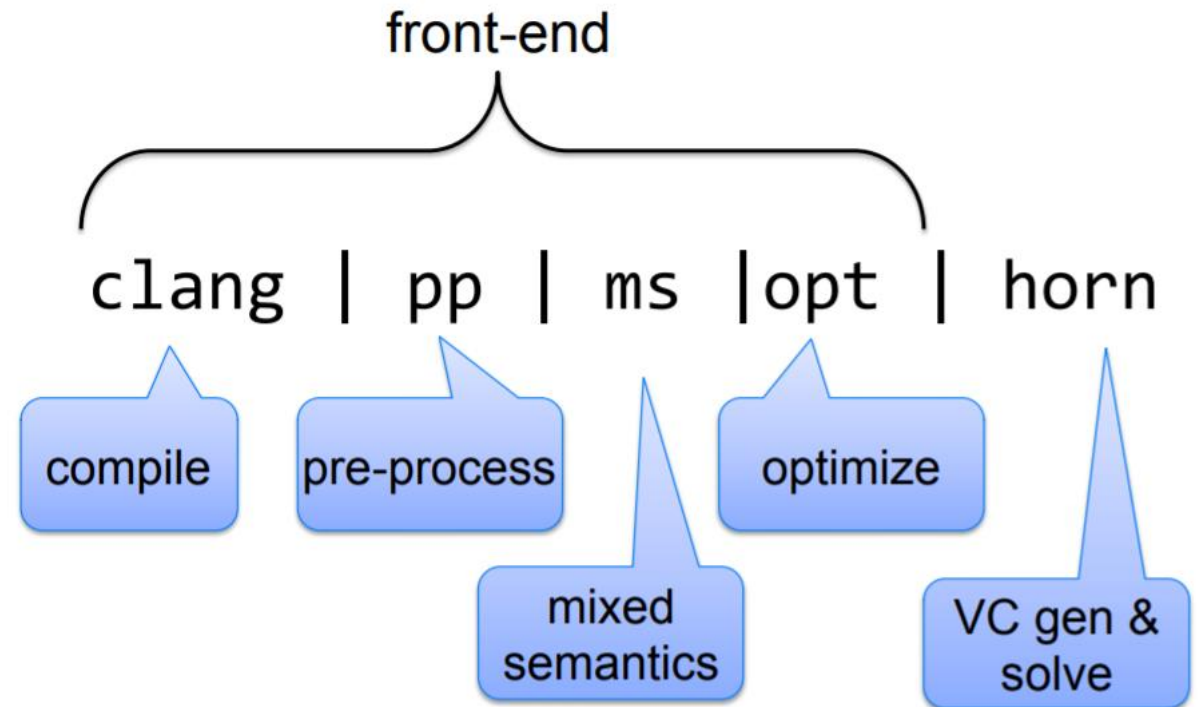
- Front-end
- Middle-end
- Back-end



## Front-end

- Takes an LLVM based program (e.g., C) input program and generated LLVM IR bytecode.
- Specifically, it performs the pre-processing and optimization of the bytecode for verification purposes.

### Verification Pipeline



## Pre-processing for Verification


- ▶ First, the input C program is pre-processed with CIL\* to insert line markings for printing user-friendly counterexamples, define missing functions that are implicitly defined, and initialized all local variables.
- ▶ Second, the result is translated into LLVM-IR bitcode, using llvm-gcc. After that, it performs compiler optimizations and preprocessing to simplify the verification task.
  - ▶ Function inlining, conversion to static single assignment (SSA) form, dead code elimination, peephole optimizations, CFG simplifications, etc.
- ▶ Make SeaHorn not to be limited to C programs, but applicable to a broader set of languages based on LLVM (C++, Objective C and Swift).

\* intermediate language and tools for analysis and transformation of C programs

# From CFG to Cut Point Graph

- ▶ A Cut Point Graph hides (summarizes) fragments of a control flow graph by (summary) edges.
- ▶ Vertices correspond to some basic blocks.
- ▶ An edge between cut-points  $c$  and  $d$  summarizes all finite (loop-free) executions from  $c$  to  $d$  that do not pass through any other cut-points.

## Cut Point Graph Example



## Mixed semantics

- ▶ One typical problem in proving safety of large programs is that assertions can be nested very deep inside the call graph. As a result, counterexamples are longer and it is harder to decide for the verification engine what is relevant for the property of interest.
- ▶ To mitigate this problem, the front-end provides a transformation based on the concept of **mixed semantics**.
  - ▶ - if P may fail, then make a copy of P's body (in main) and jump to the copy.
  - ▶ - if P may succeed, then make the call to P as usual. Since P is known not to fail each assertion in P can be safely replaced with an assume.
- ▶ Upon completion, only the main function has assertions and each procedure is inlined at most once.



## Mixed semantics - example

- ▶ A main procedure calling two other procedures  $p_1$  and  $p_2$  with three assertions  $c_1$ ,  $c_2$ , and  $c_3$ .
- ▶ The new program after the mixed-semantics transformation.



<i>main</i> () <i>p1</i> (); <i>p1</i> (); assert ( <i>c1</i> ); <i>p1</i> () <i>p2</i> (); assert ( <i>c2</i> ); <i>p2</i> () assert ( <i>c3</i> );	<i>main<sub>new</sub></i> () if (*) goto <i>p1<sub>entry</sub></i> ; else <i>p1<sub>new</sub></i> (); if (*) goto <i>p1<sub>entry</sub></i> ; else <i>p1<sub>new</sub></i> (); if ( $\neg c_1$ ) goto <i>error</i> ; assume (false);	<i>p1<sub>entry</sub></i> : if (*) goto <i>p2<sub>entry</sub></i> ; else <i>p2<sub>new</sub></i> (); if ( $\neg c_2$ ) goto <i>error</i> ; <i>p2<sub>entry</sub></i> : if ( $\neg c_3$ ) goto <i>error</i> ; assume (false); <i>error</i> : assert (false);	<i>p1<sub>new</sub></i> () <i>p2<sub>new</sub></i> (); assume ( <i>c2</i> ); <i>p2<sub>new</sub></i> () assume ( <i>c3</i> );
---	--	--	---

Fig. 2: A program before and after mixed-semantics transformation.

# Middle-end

- ▶ Encoding verification conditions:
  - ▶ SeaHorn provides two different semantics encodings: a) a small step encoding and b) a large-block encoding (LBE).
  - ▶ Different degree of precision: LBE is often more efficient but small-step might be more useful if a counterexample is needed.

```
L1:  if(p1) {
L2:      x1 = 1;
      }
L3:  if(p2) {
L4:      x2 = 2;
      }
L5:  if(p3) {
L6:      x3 = 3;
      }
L7:  if(p1) {
L8:      if (x1 != 1) goto ERR;
      }
L9:  if (p2) {
L10:     if (x2 != 2) goto ERR;
      }
L11: if (p3) {
L12:     if (x3 != 3) goto ERR;
      }
L13: return EXIT_SUCCESS;
ERR: return EXIT_FAILURE;
```



(a) Example C program

(b) ART for SBE

(c) ART for LBE

## Constrained Horn Clauses (CHC)

- ▶ A Constrained Horn Clause (CHC) is a FOL formula of the form:

$$\forall V . (\phi \wedge p_1[X_1] \wedge \dots \wedge p_n[X_n] \rightarrow h[X]),$$

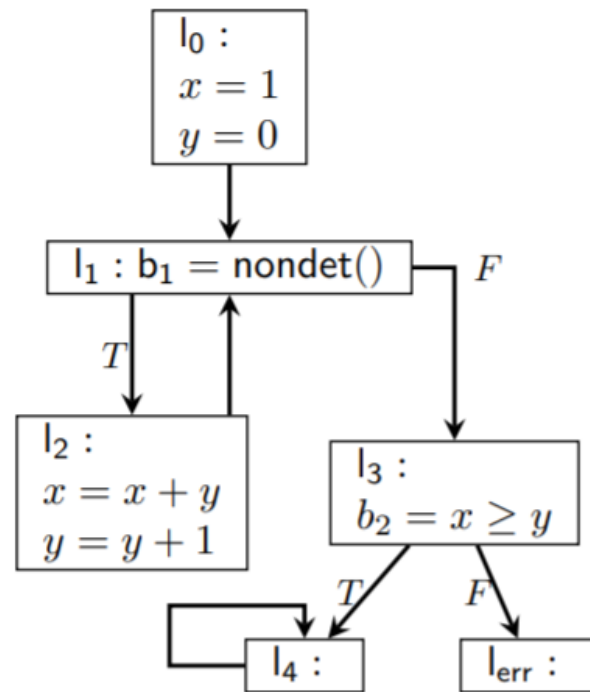
where

- $A$  is a background theory (e.g., Linear Arithmetic, Arrays, Bit-Vectors, or combinations of the above)
- $\phi$  is a constrained in the background theory  $A$
- $p_1, \dots, p_n, h$  are  $n$ -ary predicates
- $p_i[X]$  is an application of a predicate to first-order terms

## Example – small-step encoding of VCs using Horn clauses

- Program  $\rightarrow$  Control-Flow  $\rightarrow$  Verification Conditions

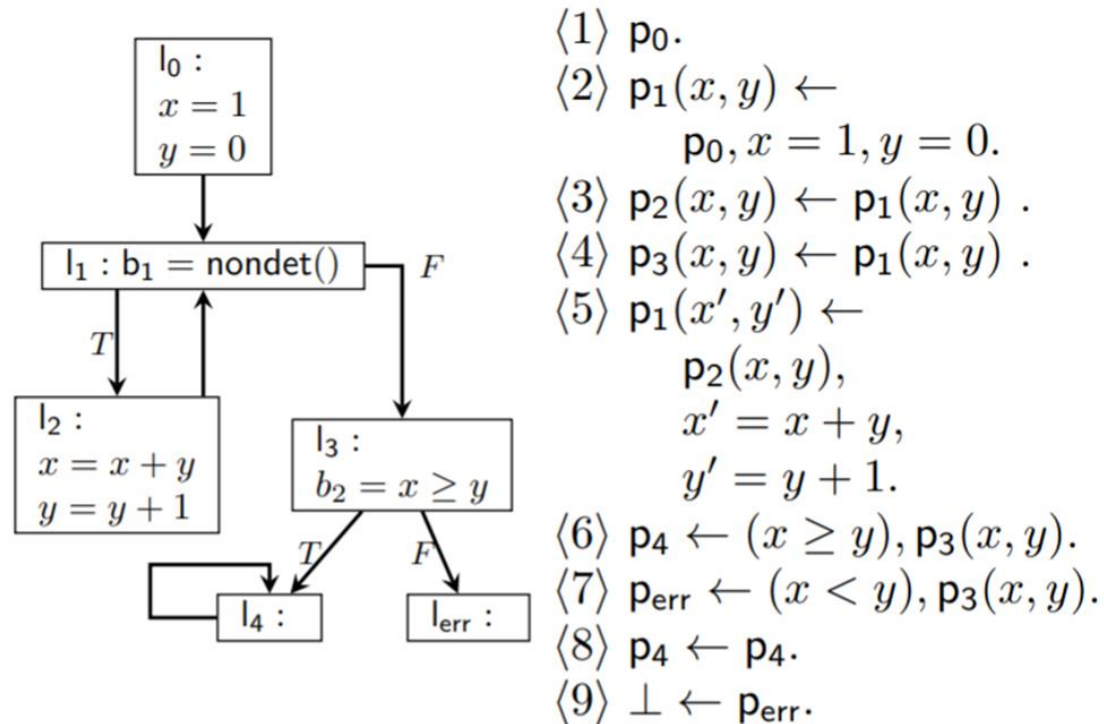
```
int x = 1;
int y = 0;
while (*) {
    x = x + y;
    y = y + 1;
}
assert(x ≥ y);
```



- $\langle 1 \rangle \text{ p}_0.$
- $\langle 2 \rangle \text{ p}_1(x, y) \leftarrow$   
 $\text{p}_0, x = 1, y = 0.$
- $\langle 3 \rangle \text{ p}_2(x, y) \leftarrow \text{p}_1(x, y).$
- $\langle 4 \rangle \text{ p}_3(x, y) \leftarrow \text{p}_1(x, y).$
- $\langle 5 \rangle \text{ p}_1(x', y') \leftarrow$   
 $\text{p}_2(x, y),$   
 $x' = x + y,$   
 $y' = y + 1.$
- $\langle 6 \rangle \text{ p}_4 \leftarrow (x \geq y), \text{p}_3(x, y).$
- $\langle 7 \rangle \text{ p}_{\text{err}} \leftarrow (x < y), \text{p}_3(x, y).$
- $\langle 8 \rangle \text{ p}_4 \leftarrow \text{p}_4.$
- $\langle 9 \rangle \perp \leftarrow \text{p}_{\text{err}}.$

## Example – small-step encoding of VCs using Horn clauses

- ▶ The set of CHCs essentially represents the small-step operational semantics of CFG.
- ▶ Each basic block is encoded as a Horn clause.
- ▶ A basic block label  $l_i$  in the CFG is translated into  $p_i(X_1, \dots, X_n)$ .




## Middle-end

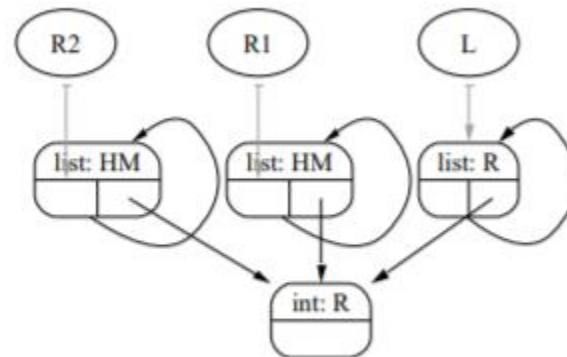
- ▶ SeaHorn middle-end offers a very simple interface for developers to implement an encoding of the verification semantics that fits their needs.
- ▶ Symbolic store:
  - the core of the SeaHorn middle-end.
  - A symbolic store simply maps program variables to symbolic values.
  - The small-step verification semantics is provided by implementing a symbolic execution interface.
- ▶ SeaHorn middle-end includes verification semantics with different levels of abstraction:
  - Registers only: only models LLVM numeric registers.
  - Registers + Pointers (without memory content): models numeric and pointer registers.
  - Registers + Pointers + Memory: models numeric and pointer registers and the heap. The heap is modeled by a collection of non-overlapping arrays.

## Middle-end


- ▶ Data structure analysis (DSA): a heap analysis to model heap.
- ▶ The analysis build for each function a DS graph where each **node** represents a potentially infinite set of memory objects and distinct DSA nodes express disjoint sets of objects. **Edges** in the graph represents points-to relationships between DS nodes.
- ▶ Given a DS graph we can map each DS node to an array. Then each memory load (read) and store (write) in the LLVM bitcode can be associated with a particular DS node.



(a) DS Graph for createnode



(b) DS Graph for splitclone



(c) DS Graph for processlist

# Verification Engines

- Horn clause-based verification tool.
- SMT-Based Model Checking with SPACER
- Abstract Interpretation with IKOS

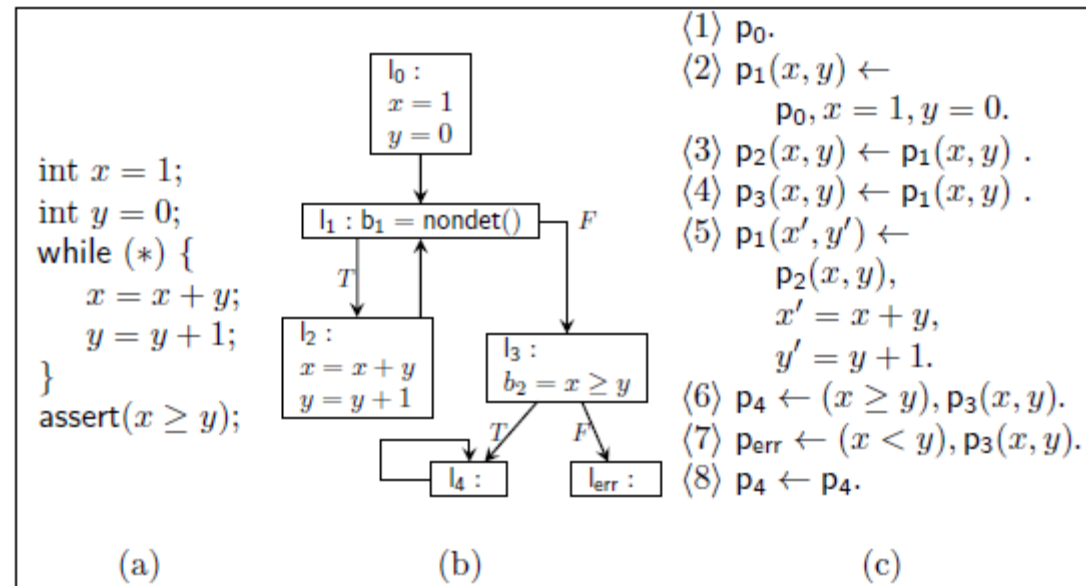


## SPACER: Solving CHC in Z3

- Spacer: solver for SMT-constrained Horn Clauses
- Support for Non-linear CHC
- Support SMT-Theories

## IKOS: an open-source library of abstract domains with a SOTA fixed-point algorithm

- ▶ SeaHorn users can choose IKOS as the only back-end engine to discharge proof obligations.
- ▶ In the example, SPACER alone can discover  $x \geq y$  but it misses the vital invariant  $y \geq 0$ . Thus, it does not terminate. On the contrary, IKOS alone with the abstract domain of DBMs(Difference-Bound Matrix) can prove safety immediately.



# Experimental Evaluation

## ▸ Results of SV-COMP 2015

### **SV-COMP 2015**

<http://sv-comp.sosy-lab.org/2015/>

4<sup>th</sup> Competition on Software Verification held (here!) at TACAS 2015

#### Goals

- Provide a snapshot of the state-of-the-art in software verification to the community.
- Increase the visibility and credits that tool developers receive.
- Establish a set of benchmarks for software verification in the community.

#### Participants:

- Over 22 participants, including most popular Software Model Checkers and Bounded Model Checkers


#### Benchmarks:

- C programs with error location (programs include pointers, structures, etc.)
- Over 6,000 files, each 2K – 100K LOC
- Linux Device Drivers, Product Lines, Regressions/Tricky examples
- <http://sv-comp.sosy-lab.org/2015/benchmarks.php>

# Experimental Evaluation

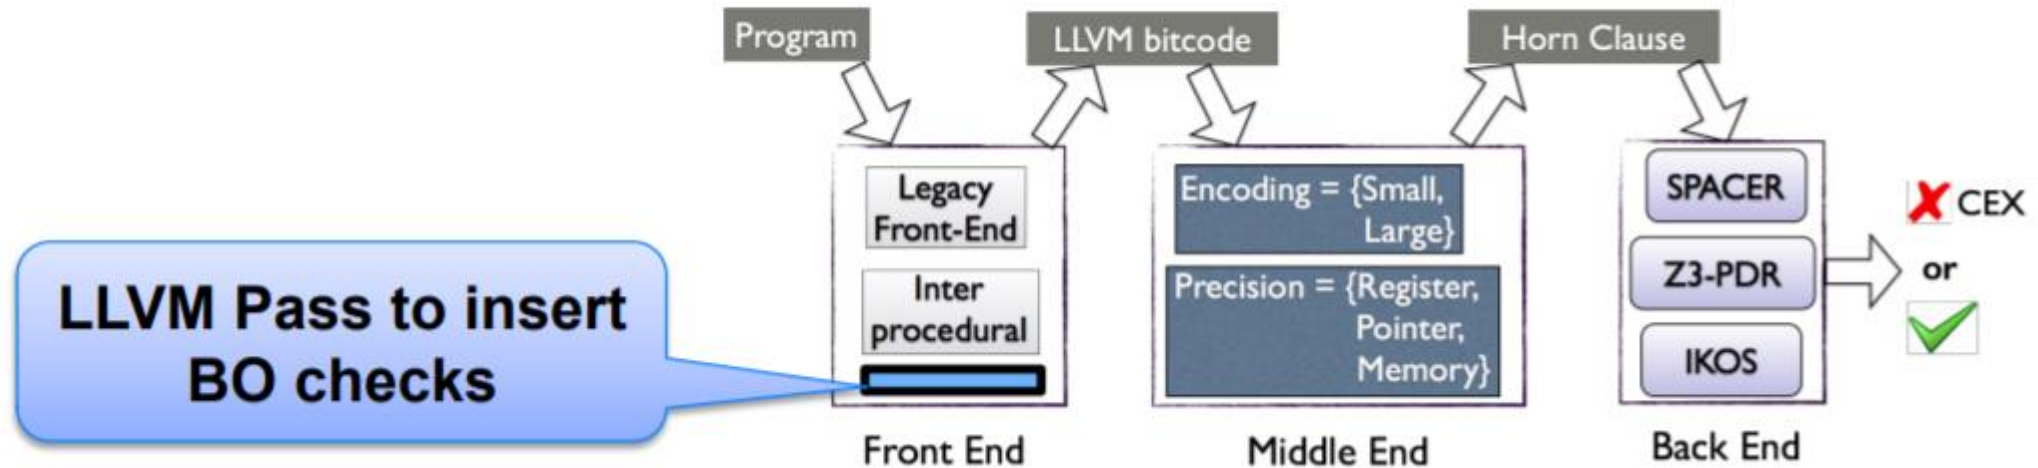
- Results of SV-COMP 2015

## Results for DeviceDriver category



## Case Study: Checking buffer overflow in Avionics Software

- Evaluated the SeaHorn built-in buffer overflow checks on two autopilot control software (paparazzi and mnav autopilots).
- To prove absence of buffer overflows, they only need to add in the front-end a new LLVM transformation pass that inserts the corresponding checks in the bitcode.



## Case Study: Checking buffer overflow in Avionics Software

- ▶ For each pointer dereference  $*p$ , add two shadow registers:  $p.offset$  and  $p.size$ .
- ▶ For each  $*p$ , add two assertions:
  - $\text{assert}(p.offset \geq 0)$  (underflow)
  - $\text{assert}(p.offset < p.size)$  (overflow)